

Linear Regression Salary by Years of Experiences

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
import warnings
warnings.filterwarnings('ignore')
import statsmodels.api as smt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
df = pd.read_csv("Salary_Data.csv")
df.head()
```

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891

```
df.shape
```

```
(30, 2)
```

```
df.describe()
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
# 0  YearsExperience  30 non-null    float64
# 1  Salary          30 non-null    float64
```

```

---  ---
0  YearsExperience  30 non-null    float64
1  Salary          30 non-null    int64
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes

df.isnull().sum()

YearsExperience    0
Salary            0
dtype: int64

for k, v in df.items():
    q1 = v.quantile(0.25)
    q3 = v.quantile(0.75)
    irq = q3 - q1
    v_col = v[(v <= q1 - 1.5 * irq) | (v >= q3 + 1.5 * irq)]
    perc = np.shape(v_col)[0] * 100.0 / np.shape(df)[0]
    print("Column %s outliers = %.2f%%" % (k, perc))

Column YearsExperience outliers = 0.00%
Column Salary outliers = 0.00%

df.skew()

YearsExperience    0.37956
Salary            0.35412
dtype: float64

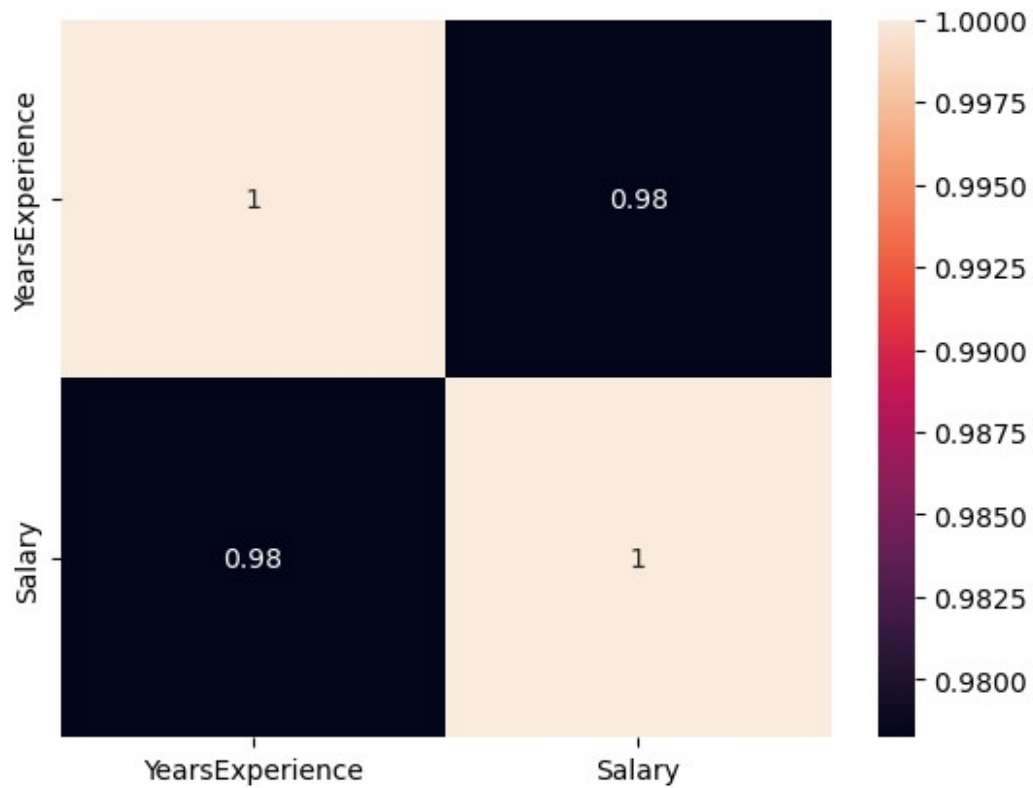
df.corr()

          YearsExperience  Salary
YearsExperience    1.000000  0.978242
Salary            0.978242  1.000000


sns.heatmap(data=df.corr(),annot=True)

<Axes: >

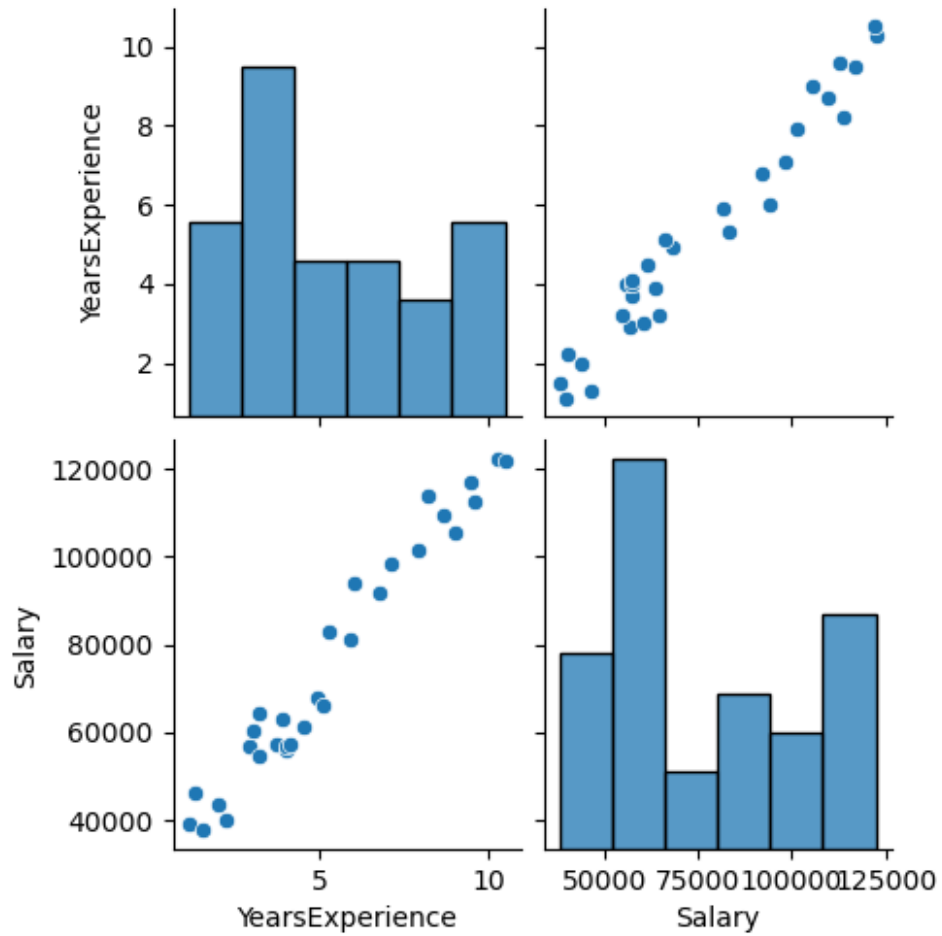
```




```
fig, axs = plt.subplots(ncols=2, nrows=1)
index = 0
axs = axs.flatten()
for k,v in df.items():
    sns.distplot(v, ax=axs[index])
    index += 1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```



```
sns.pairplot(df)  
<seaborn.axisgrid.PairGrid at 0x2c7959746d0>
```



```
sns.lmplot(x='YearsExperience', y='Salary', data=df, order=2, ci=None)
<seaborn.axisgrid.FacetGrid at 0x2c795f83390>
```



```
x =np.array(df['YearsExperience']).reshape(-1,1)
y =np.array(df['Salary']).reshape(-1,1)

x_train, x_test, y_train,y_test = train_test_split(x,y,test_size=0.30)
reg = LinearRegression()
reg.fit(x_train,y_train)
print(y_train.flatten())


[ 93940  81363  83088 116969  56642  57081  39343  98273 113812  55794
 109431  57189 112635  64445 101302  66029  63218  54445  39891  67938
 121872]

reg.intercept_
array([24648.40962547])


reg.coef_
array([[9718.66152021]])

y_predict= reg.predict(x_train)
plt.scatter(x_train,y_train,color='b')
```

```
plt.plot(x_train,y_predict,color='r')  
plt.show()
```



```
print(reg.score(x_test,y_test))  
0.9624865713263218  
  
y_predict= reg.predict(x_test)  
plt.scatter(x_test,y_test,color='b')  
plt.plot(x_test,y_predict,color='r')  
plt.show()
```



```
R2=metrics.r2_score(y_test,y_predict)
R2
0.9624865713263218
print(metrics.mean_absolute_error(y_test,y_predict))
4562.02128719395
print(metrics.mean_squared_error(y_test,y_predict))
29742511.860138074
print(np.sqrt(metrics.mean_squared_error(y_test,y_predict)))
5453.669577462323
print(reg.predict([[5]]))
[[73241.71722654]]
```


Linear Regression Assumptions

```
error= y_test-y_predict
error
```




```
array([[ -6534.3633074 ],
       [-7271.38646644],
       [-1495.4019058 ],
       [ -560.7326659 ],
       [-6566.05570633],
       [ 1002.69203707],
       [-2359.62328368],
       [ 8922.33039825],
       [ 6345.60581388]])
```


```
#Linearity
sns.regplot(x=x,y=y)
plt.show()
```



```
#Normality
sns.distplot(error)
plt.xlabel('residual')
plt.show()
```




```
#No Autocorrelation  
acf=plot_acf(error)  
plt.show()
```



```
#Homoscedasticity
print(type(error))
sns.scatterplot(x=y_predict.flatten(),y=error.flatten())
plt.xlabel('predicted values')
plt.ylabel('residuals')
plt.plot(y_predict, [0]*len(y_predict))
plt.show()

<class 'numpy.ndarray'>
```



```
#Multicollinearity  
sns.heatmap(df.corr(),annot=True)  
plt.show()
```

